

QAD

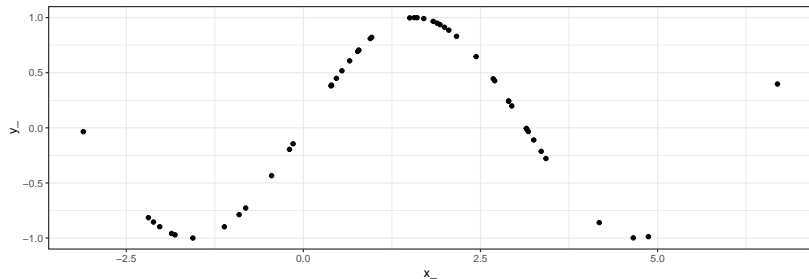
Fabian Oberreiter, Konrad Medicus, Tobias Hilgart

17.06.2019

Motivation - Asymmetric Dependence

Motivation

Bivariate sample:

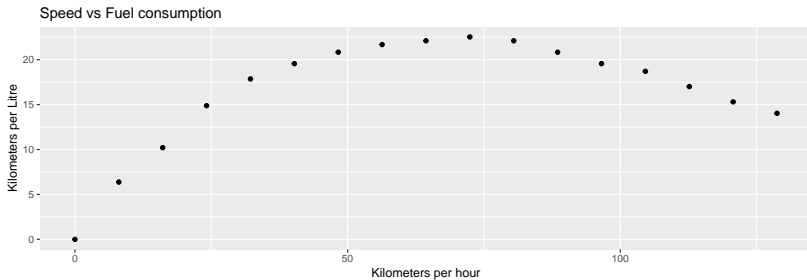


$X \sim \mathcal{N}(1.42, 2)$, $Y = \sin(X)$.

Would it be easier to predict Y given X , or X given Y ?

Motivation

Real life example: Average speed vs. Fuel consumption



Motivation

What is dependence between random variables X and Y ?

X ... result of a coin toss,

Y ... result of tossing the same coin a second time.

-> we do not gain information about Y if we know X and vice versa.

X ... result of drawing a card from a deck containing 2 cards,

Y ... result of drawing the remaining card.

-> we know everything about Y if we know X and vice versa.

Motivation

Why could this be a problem?

How would we usually quantify dependence?

Motivation

Why could this be a problem?

How would we usually quantify dependence?

Correlation coefficient:

```
cor(x_,y_)
```

```
## [1] 0.2870855
```

```
cor(y_,x_)
```

```
## [1] 0.2870855
```

Motivation

By definition of correlation (or covariance):

$$\text{corr}(X, Y) = \text{corr}(Y, X)$$

The same holds true for spearman or kendall correlation.

Correlation can NOT be used as a measure of asymmetric dependence.

Motivation

What properties should an asymmetric dependency measure q have?

- ▶ $q(X, Y) \in [0, 1]$
- ▶ $q(X, Y) = 0$ iff X and Y are independent (independence)
- ▶ $q(X, Y) = 1$ iff Y is a function of X (complete dependence)
- ▶ It may be, that $q(X, Y) \neq q(Y, X)$.
- ▶ Scale changes should not affect the outcome.

qad-measure - Construction:

- ▶ Start with a bivariate sample
- ▶ Use the pseudo-observations to construct the empirical copula E_n .
- ▶ Aggregate E_n to the empirical checkerboard copula \hat{C}_n .
- ▶ Estimate $q(A) = 3D_1(A, \Pi)$ via $q(\hat{C}_n) = 3D_1(\hat{C}_n, \Pi)$.
- ▶ For sufficiently large n , we have $q(\hat{C}_n) \approx q(A)$.

(Empirical) Copula

Idea



We have something that links univariate and bivariate distributions and contains all the information about mutual dependency:

Copula (lat. link").

It's existence is guaranteed by Sklar's Theorem, therefore it makes sense to and use it for a dependency measure.

How do we get this copula from a given bivariate sample?


Empirical copula



similarly to the empirical distribution:

We start with a sample $(x_1, y_1), \dots, (x_n, y_n)$ and the pseudo - observations (normalized ranks).

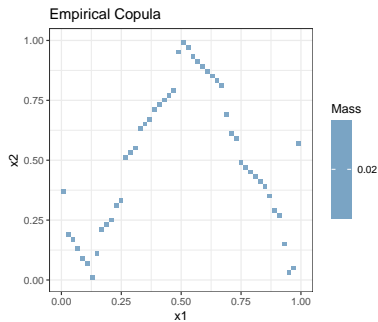
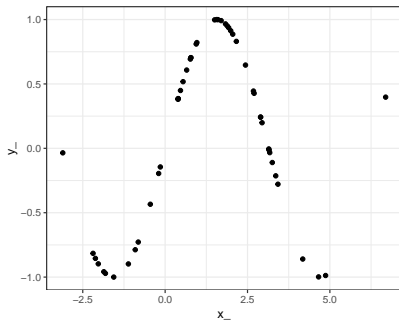
Empirical copula



And then proceed to construct an empirical copula as shown above.

Empirical copula of our example

Let's compute the empirical copula for our example:



Limitations

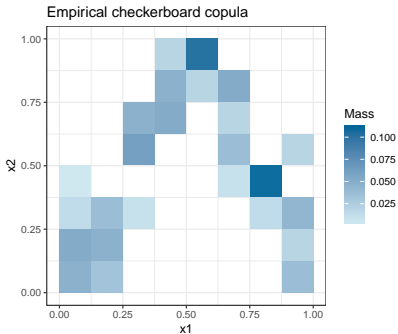
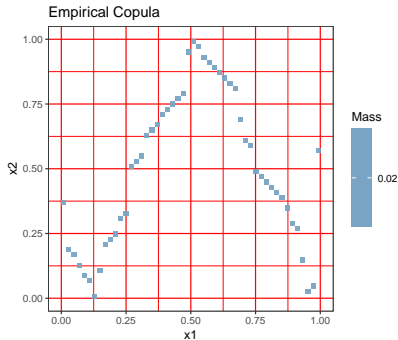
We would wish that the empirical copula E_n was a good estimate for the true underlying copula A .

Using the aforementioned metric D_1 , this is not always the case.

We need to use a different estimator for A .

Empirical checkerboard copula

We aggregate the empirical copula into what we call empirical checkerboard copula.



emp_c_copula()

The function *emp_c_copula* computes the mass-distribution of the empirical (checkerboard) copula, given a bivariate sample.

Arguments:

- ▶ *X*: a dataframe containing a bivariate sample
- ▶ *smoothing*: a logical indication whether the checkerboard copula is to be calculated
- ▶ *resolution*: an integer indicating the resolution of the checkerboard aggregation (the number of breaks in the grid)

emp_c_copula()

The function *emp_c_copula* computes the mass-distribution of the empirical (checkerboard) copula, given a bivariate sample.

```
n = 50
x = rnorm(n,0,1)
y = runif(n,0,1)
df = data.frame(x,y)
emp_cop = emp_c_copula(df,smoothing = FALSE)
emp_check_cop = emp_c_copula(df,smoothing = TRUE,resolution
```

plot_density()

The function *plot_density* plots the density/mass of the empirical (checkerboard) copula.

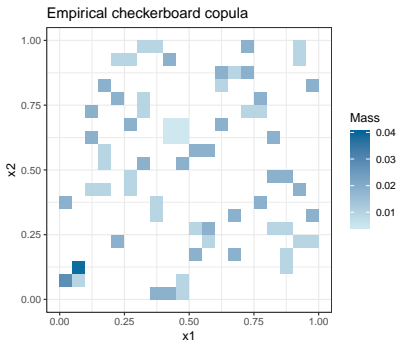
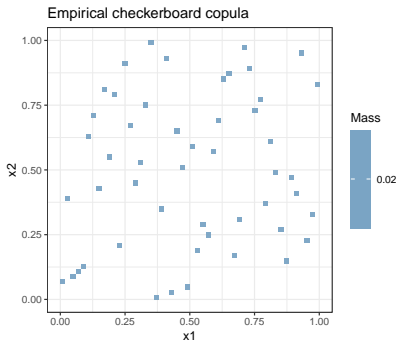
Arguments:

- ▶ **mass_matrix**: A squared matrix containing the mass distribution (output of `emp_c_copula()`)
- ▶ **density**: A logical indication whether density or mass is plotted

plot_density()

The function `plot_density` allows us to visualize copulae:

```
plot_density(emp_cop, density=FALSE)
plot_density(emp_check_cop, density=FALSE)
```



Exercise 1

Calculate the empirical copula – not the checkerboard copula – while having the argument *smoothing*=*TRUE*. (Hint: Resolution parameter)

Optional: Visualize this by plotting the pseudo observations of your generated sample into the mass-plot (Compare: slide 13,14).

Explain the differences in the following plots:

```
n=10
df = data.frame(x = rnorm(n), y = runif(n))
emp_cop = emp_c_copula(df, smoothing = FALSE)
emp_check_cop = emp_c_copula(df, smoothing = TRUE, resolution = 10)
plot_density(emp_cop, density = FALSE)
plot_density(emp_check_cop, density = FALSE)
```

The dependency measure

The dependency measure

Based on a metric D_1 for copulae, a dependency measure can be constructed in the following way:

$$q(A) := 3 \cdot D_1(A, \Pi) \in [0, 1]$$

Π denotes the product copula, which stems from two completely independent RVs.

In a way, we measure the “distance” to complete independence.

qad()

The function *qad* quantifies the asymmetric dependence of two RVs X and Y . It achieves this by calculating the empirical checkerboard copula to estimate the dependency measure q .

Arguments:

- ▶ X : a dataframe containing a bivariate sample in two columns.
- ▶ resolution: resolution of the emp. copula
- ▶ premutation: a logical indicating, whether a permuted p -value is computed.

qad()

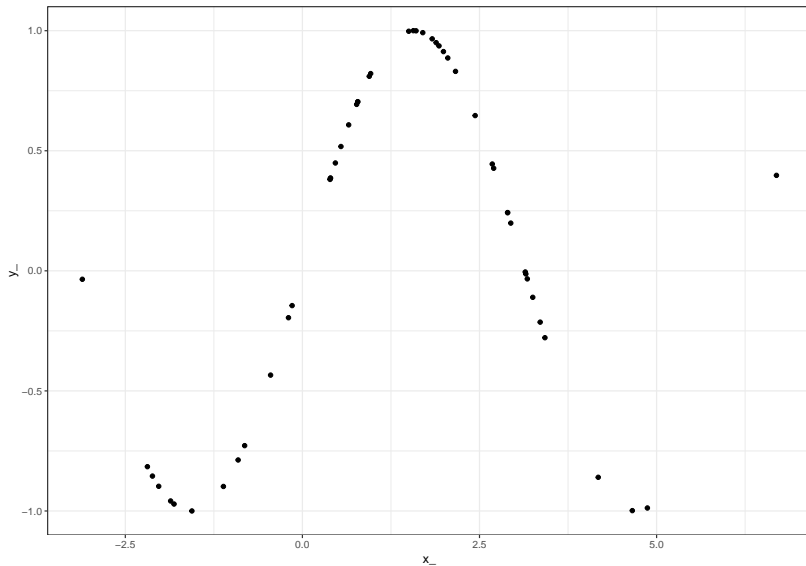
The function *qad* quantifies the asymmetric dependence of two RVs X and Y . It achieves this by calculating the empirical checkerboard copula to estimate the dependency measure q .

There are multiple ways to see the results of the *qad* calculation:

- ▶ `summary(qad(X))`
- ▶ `coef(qad(X))`
- ▶ `qad(X)$results`

qad Example 1

Back to our initial example:



qad Example 1

We take a look at the mutual dependencies.

```
mod = qad(X, print = FALSE)
coef(mod, select = c("q(x1,x2)", "q(x2,x1)", "mean.dependence",
                     "asymmetry"))
```

##	q(x1,x2)	q(x2,x1)	mean.dependence	as
##	0.7885979	0.4193818	0.6039898	0.

We see that X predicts Y better than Y predicts X .

qad Example 2

```
n = 200
x2 = runif(n, -2, 4)
y2 = exp(x2)
df = data.frame(x2, y2)
model = qad(df, print = FALSE, permutation = TRUE)
model$results
```

##		coef	p.values
## 1	q(x1,x2)	0.9608946	0
## 2	q(x2,x1)	0.9608946	0
## 3	mean.dependence	0.9608946	0
## 4	asymmetry	0.0000000	1

pairwise.qad

With the function *pairwise.qad*, we can apply the *qad* function on each pair of columns of the given dataframe. Other arguments are the same as for *qad*.

```
n = 1000
x = runif(n,0,1)
y = x^2+rnorm(n,0,1)
z = runif(n,0,1)
df = data.frame(x,y,z)


model = pairwise.qad(df, permutation = TRUE)

## [1] "Process..."
## [1] "Process: 1 / 3"
## [1] "Process: 2 / 3"
```

heatmap.qad

We can plot the results as a heatmap:

```
heatmap.qad(model, select = "dependence", fontsize = 8,  
            significance = TRUE)
```



heatmap.qad

Arguments:

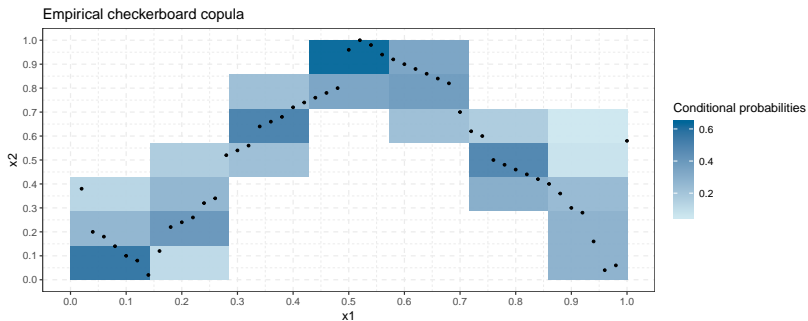
- ▶ `pw_qad`: the output of the function *pairwise.qad()*
- ▶ `select`: dependence value to be plotted: “dependence”, “mean.dependence” or “asymmetry”
- ▶ `significance`: logical; if p-values were calculated, marks the significant dependency values with a star
- ▶ `sign.level`: manually set significance level (default: 0.05)

Exercise 2

- ▶ Download the RTR-dataset
(via `load(url("http://www.trutschnig.net/RTR.RData"))`)
and sample 1000 observations.
- ▶ Examine which columns could have interesting dependencies.
- ▶ Why does this not make sense for all columns?
- ▶ Visualize the dependencies for the relevant columns. Justify why you chose them.

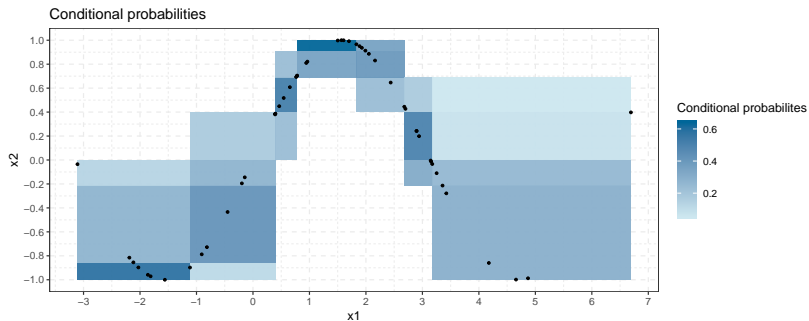
plot.qad

```
plot(mod, addSample = TRUE, copula = TRUE)
```



plot.qad

```
plot(mod, addSample = TRUE)
```



predict.qad

The function *predict.qad()* predicts the probabilities to end up in specific intervals given x or y values.

The prediction can be computed in the copula or the retransformed data setting.

```
predict(mod, values = 1.42, conditioned = "x1", nr_intervals = 5)
```

```
##           (-1,-0.79] (-0.79,-0.02] (-0.02,0.43] (0.43,0.82] (0.82,1.00]  
## 1.42                0                0                0                0.216
```

Exercise 3

- ▶ Create a dummy dataset: $n = 100$ observations of a uniformly distributed random variable and the corresponding square values.
- ▶ Create a new qad object and use the plot function to visualize the copula.
- ▶ Add the observations to the plot with the addSample parameter.
- ▶ Use the predict.qad method and sum up the probabilities for one strip (horizontal or vertical) for one value.

cci - Conditional Confidence Interval

We can also use the qad metric to test hypothesis:

cci provides a confidence interval for independence (for the qad measure).

We can compare the qad-value to the interval boundaries and check if the null hypothesis of independence holds or has to be rejected.

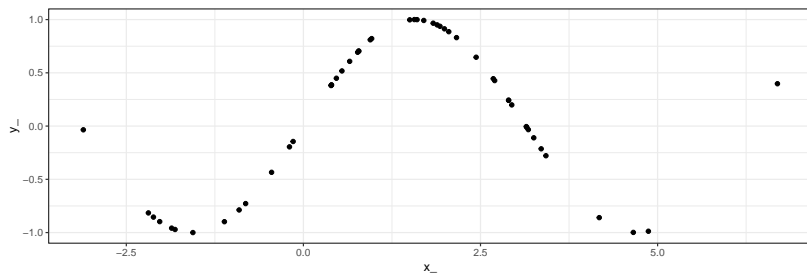
cci

We calculate our boundaries

```
c = cci(n, alternative = "one.sided")
```

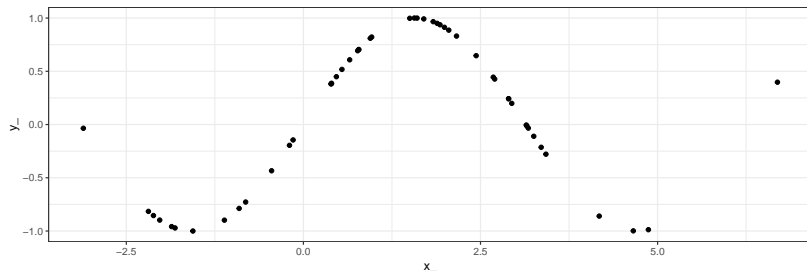
and check, whether the calculated dependence for our example lies between them.

```
if(coef(mod, select = 'q(x1,x2)') %in% c){  
  print('Accept H0')  
}else{  
  print('Reject H0')  
}
```



And indeed, for our example, we reject the hypothesis of independence:

```
## [1] "q(x1,x2) not in [ 0 , 0.1813481327652 ]"  
## [1] "Reject H0"
```

Similarly for the other direction:

```
## [1] "q(x2,x1) not in [ 0 , 0.1813481327652 ]"  
## [1] "Reject H0"
```

Final Exercise 4:

Take a look at the *attitude* dataset:

Find out the two attributes with the highest (one sided) dependency.
Provide respective plots.

Are those dependencies symmetric? Argue by providing p-values for asymmetry (hint: do not to calculate the p-values for all pairs, only the chosen two)

Reject or Accept the hypothesis of independence for your chosen pairs.